

▼ Problem Set 2: CNNs & TensorFlow [75 pts]

Note: The following has been verified to work with TensorFlow 2.0

* Adapted from official TensorFlow™ tour guide.

TensorFlow is a powerful library for doing large-scale numerical computation. One of the tasks at which it excels is implementing and training deep neural networks. In this assignment you will learn the basic building blocks of a TensorFlow model while constructing a deep convolutional MNIST classifier.

What you are expected to implement in this tutorial:

- Create a softmax regression function that is a model for recognizing MNIST digits, based on looking at every pixel in the image
- Use Tensorflow to train the model to recognize digits by having it "look" at thousands of examples
- Check the model's accuracy with MNIST test data
- Build, train, and test a multilayer convolutional neural network to improve the results

▼ Part 1: Coding [50 pts]

▼ Data

After importing tensorflow, we can download the MNIST dataset with the built-in TensorFlow/Keras method.

```
import os

import tensorflow as tf
import matplotlib.pyplot as plt

os.environ['OMP_NUM_THREADS'] = '1'
tf.__version__

'2.6.0'

(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data(
    class_names = ['0', '1', '2', '3', '4',
```


'5', '6', '7', '8', '9']

```
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[train_labels[i]])
plt.show()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist>

11493376/11490434 [=====] - 0s 0us/step

11501568/11490434 [=====] - 0s 0us/step



▼ Build the CNN

In this part we will build a customized TF2 Keras model. As input, a CNN takes tensors of shape (image_height, image_width, color_channels), ignoring the batch size. For MNIST, you will configure our CNN to process inputs of shape (28, 28, 1), which is the format of MNIST images. You can do this by passing the argument input_shape to our first layer.

The overall architecture should be:

Model: "customized_cnn"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	multiple	320

max_pooling2d_1 (MaxPooling2D)	multiple	0

conv2d_3 (Conv2D)	multiple	18496

flatten_1 (Flatten)	multiple	0

dense_2 (Dense)	multiple	7930880

dense_3 (Dense)	multiple	10250
=====		
Total params: 7,959,946		
Trainable params: 7,959,946		
Non-trainable params: 0		

First Convolutional Layer [5 pts]

We can now implement our first layer. The convolution will compute 32 features for each 3x3 patch. The first two dimensions are the patch size, the next is the number of input channels, and the last is the number of output channels.

Max Pooling Layer [5 pts]

We stack max pooling layer after the first convolutional layer. These pooling layers will perform max pooling for each 2x2 patch.

Second Convolutional Layer [5 pts]

In order to build a deep network, we stack several layers of this type. The second layer will have 64 features for each 3x3 patch.

Fully Connected Layers [10 pts]

Now that the image size has been reduced to 11x11, we add a fully-connected layer with 128 neurons to allow processing on the entire image. We reshape the tensor from the second convolutional layer into a batch of vectors before the fully connected layer.

The output layer should also be implemented via a fully connect layer.

Complete the Computation Graph [15 pts]

Please complete the following function:

```
def call(self, inputs, training=None, mask=None):
```

To apply the layer, we first reshape the input to a 4d tensor, with the second and third dimensions corresponding to image width and height, and the final dimension corresponding to the number of color channels (which is 1).

We then convolve the reshaped input with the first convolutional layer and then the max pooling followed by the second convolutional layer. These convolutional layers and the pooling layer will reduce the image size to 11x11.

Dropout Layer [5 pts]

Please add dropouts during training before each fully connected layers, as this helps avoid overfitting during training. <https://www.cs.toronto.edu/~hinton/absps/JMLRdropout.pdf>

```
class CustomizedCNN(tf.keras.models.Model):
```

```
    def __init__(self, *args, **kwargs):
        super().__init__()
        self.layer1 = tf.keras.layers.Conv2D(32,(3,3),padding='valid', activation='relu', inp
        self.maxpooling = tf.keras.layers.MaxPooling2D(pool_size=(2,2))
        self.layer2 = tf.keras.layers.Conv2D(64,(3,3),padding='valid', activation='relu', inpu
        self.flatten = tf.keras.layers.Flatten()
        self.dropout1 = tf.keras.layers.Dropout(.06)
        self.densedconnect = tf.keras.layers.Dense(128, activation='relu')
        self.dropout2 = tf.keras.layers.Dropout(.06)
        self.denseoutput = tf.keras.layers.Dense(10)
```

```
        #raise NotImplementedError('Implement Using Keras Layers.')
```

```
    def call(self, inputs, training=None, mask=None):
        inputs = tf.reshape(inputs, [-1,28,28,1] )
        inputs = self.layer1(inputs)
        inputs = self.maxpooling(inputs)
        inputs = self.layer2(inputs)
        inputs = self.flatten(inputs)
        inputs = self.dropout1(inputs)
        inputs = self.densedconnect(inputs)
```

```

inputs = self.dropout2(inputs)
inputs = self.denseoutput(inputs)
return inputs
#raise NotImplementedError('Build the CNN here.')

```

▼ Build the Model

```

model = CustomizedCNN()
model.build(input_shape=(None, 28, 28))
model.summary()

```

Model: "customized_cnn"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	multiple	320
max_pooling2d (MaxPooling2D)	multiple	0
conv2d_1 (Conv2D)	multiple	18496
flatten (Flatten)	multiple	0
dropout (Dropout)	multiple	0
dense (Dense)	multiple	991360
dropout_1 (Dropout)	multiple	0
dense_1 (Dense)	multiple	1290
=====		
Total params: 1,011,466		
Trainable params: 1,011,466		
Non-trainable params: 0		

We can specify a loss function just as easily. Loss indicates how bad the model's prediction was on a single example; we try to minimize that while training across all the examples. Here, our loss function is the cross-entropy between the target and the softmax activation function applied to the model's prediction. As in the beginners tutorial, we use the stable formulation:

```

model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

```

▼ Train and Evaluate the Model [5 pts]

We will use a more sophisticated ADAM optimizer instead of a Gradient Descent Optimizer.

Feel free to run this code. Be aware that it does 10 training epochs and may take a while (possibly up to half an hour), depending on your processor.

The final test set accuracy after running this code should be approximately 98.7% -- not state of the art, but respectable.

We have learned how to quickly and easily build, train, and evaluate a fairly sophisticated deep learning model using TensorFlow.

```
train_image = tf.cast(train_images, tf.float32)
validate = tf.cast(test_images, tf.float32)
history = model.fit(train_image, train_labels, epochs = 10, batch_size = 1, validation_data =
```

```
Epoch 1/10
60000/60000 [=====] - 243s 4ms/step - loss: 0.2681 - accuracy:
Epoch 2/10
60000/60000 [=====] - 246s 4ms/step - loss: 0.1930 - accuracy:
Epoch 3/10
60000/60000 [=====] - 239s 4ms/step - loss: 0.1802 - accuracy:
Epoch 4/10
60000/60000 [=====] - 237s 4ms/step - loss: 0.1581 - accuracy:
Epoch 5/10
60000/60000 [=====] - 235s 4ms/step - loss: 0.1480 - accuracy:
Epoch 6/10
60000/60000 [=====] - 239s 4ms/step - loss: 0.1549 - accuracy:
Epoch 7/10
60000/60000 [=====] - 243s 4ms/step - loss: 0.1511 - accuracy:
Epoch 8/10
60000/60000 [=====] - 241s 4ms/step - loss: 0.1441 - accuracy:
Epoch 9/10
60000/60000 [=====] - 241s 4ms/step - loss: 0.1397 - accuracy:
Epoch 10/10
60000/60000 [=====] - 241s 4ms/step - loss: 0.1425 - accuracy:
```

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'validation accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.9, 1])
plt.legend(loc='lower right')
```

```
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(test_acc)
```

▼ Part 2: Written [25 pts]

2.1 [5pts] Suppose we have a neural network with ReLU activation function. Let's say, we replace ReLU activations by linear activations. Would this new neural network be able to approximate a non-linear function? And why?

Solution Here No because the ReLU function slopes while a non-linear function would not slope and choose either zero and one and not have data for the inbetween. You would need a non-linear function to approximate a non-linear function.

2.2 [5pts] List two ways to downsize feature maps in convolutional neural networks.

Solution Here Max Pooling and
ReLU function

2.3 [5pts] Describe K-fold cross-validation.

Solution Here Divide dataset into K folds and then each fold is used as a testing set. Ensures that every observation from the original dataset has the chance of appearing in training and test set.

2.4 [5pts] List four techniques that help a model avoid overfitting. Briefly explain each in 2-3 sentences.

Solution Here Cross Validation is used to split the data into multiple sets, the test set is unforeseen data.

Early stopping is used to early stop the training data, makes sure the validation set doesn't increase. This helps make sure that you don't overtrain your data. There's a point where the training is helpful, and a point where the training is no longer helpful to get your solution to data.

Dropout layers randomly drops connections and set it to zero, model averaging.

Data augmentation increases the data, minorly editing data and adding it to existing data which reduces model variance.

2.5 [5pts] Read the following on convolutions: <https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1>. What is padding? And what does it help us achieve.

Solution Here Padding extends the area of which a convolutional neural network processes an image. Allows more space for the filter to cover in the image. It is like a magnifying glass, giving more data into a smaller area. It makes the result more accurate when it comes to images.

✓ 40m 22s completed at 8:37 PM

